

Student: _____

Class: _____ No.: _____

1st test

Duration: **75 min.**

Date: 12 de april 2025

Carefully read the information on this page while you wait for the lecturer's instruction to begin the exam.

This exam contains 7 page(s), including **one** cover sheet (this one), **four** question and answer sheets, and **one** blank sheet. The exam consists of **6 questions**, some of which are divided into sub-questions, making a total of **20 points**.

Recommendations:

- It is advised that you read all the exam questions before starting, in order to better **plan your answering strategy**;
- True / False questions may have zero, one, or several correct answers;
- For True / False items, marks will only be awarded for answers that clearly indicate the intended option;
- In True / False questions:
 - **each incorrect answer deducts $\frac{1}{4}$ of the question's value**. You may choose not to answer, in which case no marks will be awarded or deducted;
 - Answers you wish to mark as **correct** should be **circled**, e.g. True False. **Illegible** answers will be marked as incorrect;
 - To correct an answer, **and only if you cannot erase it**, mark the incorrect answer with a **clearly visible cross** and then mark the correct one (True or False).
- You may use pencil and eraser, but do not forget to finalise your answers in **black ink** (preferably);
- For open-ended questions, justify and support your answers appropriately. The **simplicity** and **clarity** of your answers will be taken into account in the assessment;
- If you consider any **question ambiguous**, explain the interpretation you used when answering. **Your interpretation of the question is part of the assessment**;

By submitting the test, you agree with the conditions stated above.

Table (for the **EXCLUSIVE** use of the lecturer)

Question:	1	2	3	4	5	6	Total
Value:	3	3	4	4	3	3	20
Grade:							

1. (3 pontos) Considering the acronym ACID, which represents the desirable properties in transactional processing systems, including database management systems (DBMS), indicate what the letter I in the acronym stands for and relate it to the *Two-Phase Locking* (2PL) protocol.

[illegible]

2. (3 pontos) The properties of *cascadeless* and *strict* schedules are similar because both:
- (a) True False always allow conflict-serialisable schedules.
 - (b) True False do not allow a transaction to write to a data item that has been written by another transaction which has not yet completed.
 - (c) True False do not allow *Write/Write* conflicts.
 - (d) True False protect against *dirty reads*.

3. Consider the *phantom records* anomaly, which results from the concurrent execution of transactions.
- (a) (1 ponto) Describe the anomaly, providing an example of a schedule that illustrates your explanation.

[illegible]

(b) (3 pontos) Considering the SQL instructions INSERT and UPDATE, indicate under what conditions they may cause the anomaly, assuming they are executed on a database that implements the 2PL protocol.

4. Consider the transactional processes below, executed concurrently on the table `conta(id integer primary key, saldo real);`. Also assume that BL-xx represent blocks of instructions executed sequentially without interruption from other transactions.

T1

T2

1

-- BI-1A:

2

begin transaction;

3

set transaction isolation level repeatable read;

4

5

-- BI-1B:

6

select * from conta where id = 1111;

7

8

-- BI-1C:

9

select * from conta where id = 2222;

10

commit;

1

-- BI-2A:

2

update conta set saldo = saldo + 500 where id =

3

1111;

4

-- BI-2B:

5

update conta set saldo = saldo + 500 where id =

6

2222;

(a) (2 pontos) Considering the schedule <BL-1A, BL-2A, BL-1B, BL-2B, BL-1C>, indicate at the end of each block of instructions the number and type of locks, assuming that the DBMS implements the 2PL protocol.

- (b) (2 pontos) For each of the transactional processes T1 and T2, indicate the possible anomalies that may occur, providing a schedule that is susceptible to those anomalies.

5. (3 pontos) Considering the `account` table defined below, write a function in `plpgsql` that uses a cursor to return all rows from the `account` table where the `balance` field is greater than a value passed as a parameter. You should take into account that: (a) Each row is returned using `RETURN NEXT`; (b) The function must return `SETOF account`.

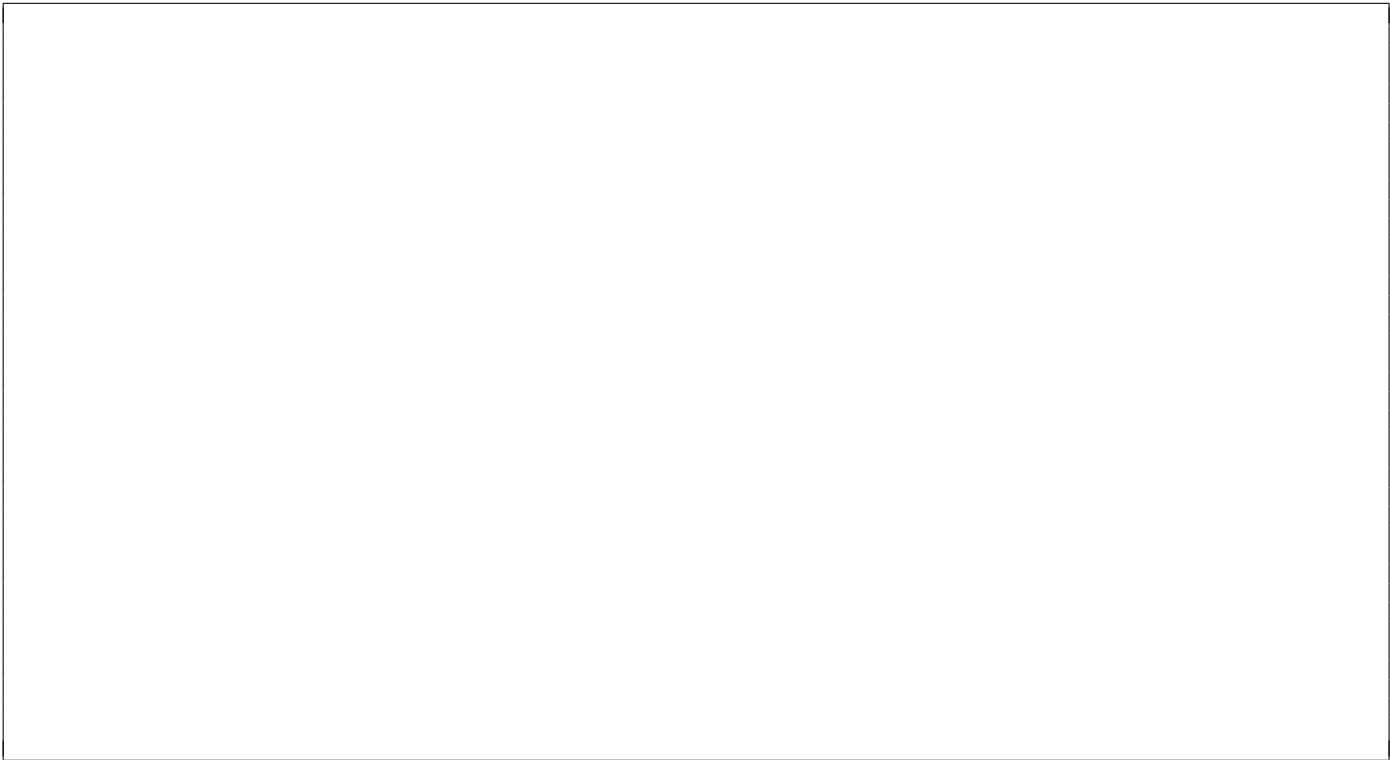
```
1 create table public.account
2 (
3     id serial primary key,
4     "description" varchar(100) not null,
5     balance money not null,
6     client int not null references client,
7     updated timestamp with time zone null
8 );
```

```
1 CREATE FUNCTION get_accounts(v money)
2 RETURNS SETOF account AS
3 $$
4 -- Implementar
5 %$$ LANGUAGE plpgsql;
```

6. (3 pontos) Write a stored procedure in PostgreSQL that uses the `get_accounts_balance` function and processes the results internally:
- (a) The procedure must accept one input parameter: a minimum balance (of type money);
 - (b) It must internally invoke the function `get_accounts_balance(v money)`;
 - (c) It must iterate over the returned rows using a FOR loop;
 - (d) It must use `raise notice` to display the values of the `description` and `balance` fields for each tuple.

Tip: You may use a variable of type ROWTYPE to store each row.

```
1 CREATE OR REPLACE PROCEDURE accounts_balance(min_balance money)
2 LANGUAGE plpgsql
3 -- Implementar
```



FIM

FOLHA EM BRANCO